

CISB 60 – ML and DL (Fall, 2024)

Final Project: Predicting FIFA Player Potential

```
In [2]: # Edit all the Markdown cells below with the appropriate information
# Run all cells, containing your code
# Save this Jupyter with the outputs of your executed cells
# PS: Save again the notebook with this outcome.
# PSPS: Don't forget to include the dataset in your submission
```

Team:

- Mohammed Khan

Course: CISB 60 – ML and DL (Fall, 2024)

Problem Statement

- This project is about house price predictions.
- **Keywords:** House price prediction, real estate ,...,

```
In [ ]: ### **Project Description**

**Objective:**
This project aims to predict FIFA player potential using machine learning and various attributes, the models can help identify high-potential players for scouting and recruitment.

**Dataset Description:**
- **Source:** FIFA dataset containing 51 attributes of players.
- **Key Features:** Includes physical characteristics, skill metrics, and overall rating.
- **Structure:**
  - Total Records: 17,954
  - Columns: Player name, age, overall rating, potential, and more.

**Business Problem:**
Football clubs need efficient ways to identify promising players. By leveraging machine learning, clubs can make better decisions while reducing scouting risks.

**Keywords:** FIFA, player potential prediction, machine learning, deep learning
```

```
In [ ]: ### **Problem Statement**

This project aims to predict FIFA player potential using advanced machine learning models. The dataset contains various attributes of players, such as physical, skill, and overall rating. By analyzing these attributes, we aim to build models that assist in identifying high-potential players.
```

```
In [ ]: ### **Keywords:**  
FIFA, player potential prediction, machine learning, deep learning, football ar
```

Required packages

- Add instructions to install the required packages

```
In [3]: ## Your code begins here
```

Methodology

1. Explain your ML and DL methodology ML (Machine Learning) Methodology: In the machine learning section, we aim to predict player potential using the K-Nearest Neighbors (KNN) algorithm. KNN is a simple yet effective algorithm that classifies data points or makes predictions based on the proximity of data points in feature space. For this project: Data is first cleaned, scaled, and prepared. KNN identifies players with similar attributes to predict their potential. Evaluation metrics like RMSE and score assess model accuracy. DL (Deep Learning) Methodology: Deep learning involves using artificial neural networks to model complex, non-linear relationships in the data. In this project: A neural network is built with multiple layers (input, hidden, output) to predict player potential. The model is trained using backpropagation and optimized with the Adam optimizer. Metrics like Mean Absolute Error (MAE) and loss curves assess performance.
2. Introduce the topics you used in your project

- Model 1
 - KNN Description:

KNN is chosen for its simplicity and interpretability. It predicts a player's potential by considering the average potential of their nearest neighbors in feature space. The number of neighbors (k) is optimized through hyperparameter tuning. Distance metrics (e.g., Euclidean) determine similarity.

- Model 2
 - Deep Learning Description:

The DNN consists of multiple fully connected layers with ReLU activation functions. Dropout layers prevent overfitting by randomly deactivating neurons during training. The final output layer uses a linear activation function to predict player potential. Metrics such as training and validation loss are visualized to track learning progress.

Your code starts here

```
In [1]: # Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
```

```
In [2]: # Load the dataset
fifa_data = pd.read_csv("fifa_players.csv")
```

```
In [3]: # Display dataset information  
print("Dataset Overview:")  
fifa_data.info()
```

Dataset Overview:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 17954 entries, 0 to 17953

Data columns (total 51 columns):

#	Column	Non-Null Count	Dtype
0	name	17954 non-null	object
1	full_name	17954 non-null	object
2	birth_date	17954 non-null	object
3	age	17954 non-null	int64
4	height_cm	17954 non-null	float64
5	weight_kgs	17954 non-null	float64
6	positions	17954 non-null	object
7	nationality	17954 non-null	object
8	overall_rating	17954 non-null	int64
9	potential	17954 non-null	int64
10	value_euro	17699 non-null	float64
11	wage_euro	17708 non-null	float64
12	preferred_foot	17954 non-null	object
13	international_reputation(1-5)	17954 non-null	int64
14	weak_foot(1-5)	17954 non-null	int64
15	skill_moves(1-5)	17954 non-null	int64
16	body_type	17954 non-null	object
17	release_clause_euro	16117 non-null	float64
18	national_team	857 non-null	object
19	national_rating	857 non-null	float64
20	national_team_position	857 non-null	object
21	national_jersey_number	857 non-null	float64
22	crossing	17954 non-null	int64
23	finishing	17954 non-null	int64
24	heading_accuracy	17954 non-null	int64
25	short_passing	17954 non-null	int64
26	volleys	17954 non-null	int64
27	dribbling	17954 non-null	int64
28	curve	17954 non-null	int64
29	freekick_accuracy	17954 non-null	int64
30	long_passing	17954 non-null	int64
31	ball_control	17954 non-null	int64
32	acceleration	17954 non-null	int64
33	sprint_speed	17954 non-null	int64
34	agility	17954 non-null	int64
35	reactions	17954 non-null	int64
36	balance	17954 non-null	int64
37	shot_power	17954 non-null	int64
38	jumping	17954 non-null	int64
39	stamina	17954 non-null	int64
40	strength	17954 non-null	int64
41	long_shots	17954 non-null	int64
42	aggression	17954 non-null	int64
43	interceptions	17954 non-null	int64
44	positioning	17954 non-null	int64
45	vision	17954 non-null	int64
46	penalties	17954 non-null	int64
47	composure	17954 non-null	int64
48	marking	17954 non-null	int64
49	standing_tackle	17954 non-null	int64
50	sliding_tackle	17954 non-null	int64

dtypes: float64(7), int64(35), object(9)
memory usage: 7.0+ MB

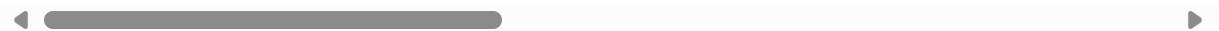
```
In [4]: # Display dataset information
print("Dataset first 5 information:")
fifa_data.head()
```

Dataset first 5 information:

Out[4]:

	name	full_name	birth_date	age	height_cm	weight_kgs	positions	nationality	overall
0	L. Messi	Lionel Andrés Messi Cuccittini	6/24/1987	31	170.18	72.1	CF,RW,ST	Argentina	
1	C. Eriksen	Christian Dannemann Eriksen	2/14/1992	27	154.94	76.2	CAM,RM,CM	Denmark	
2	P. Pogba	Paul Pogba	3/15/1993	25	190.50	83.9	CM,CAM	France	
3	L. Insigne	Lorenzo Insigne	6/4/1991	27	162.56	59.0	LW,ST	Italy	
4	K. Koulibaly	Kalidou Koulibaly	6/20/1991	27	187.96	88.9	CB	Senegal	

5 rows × 51 columns



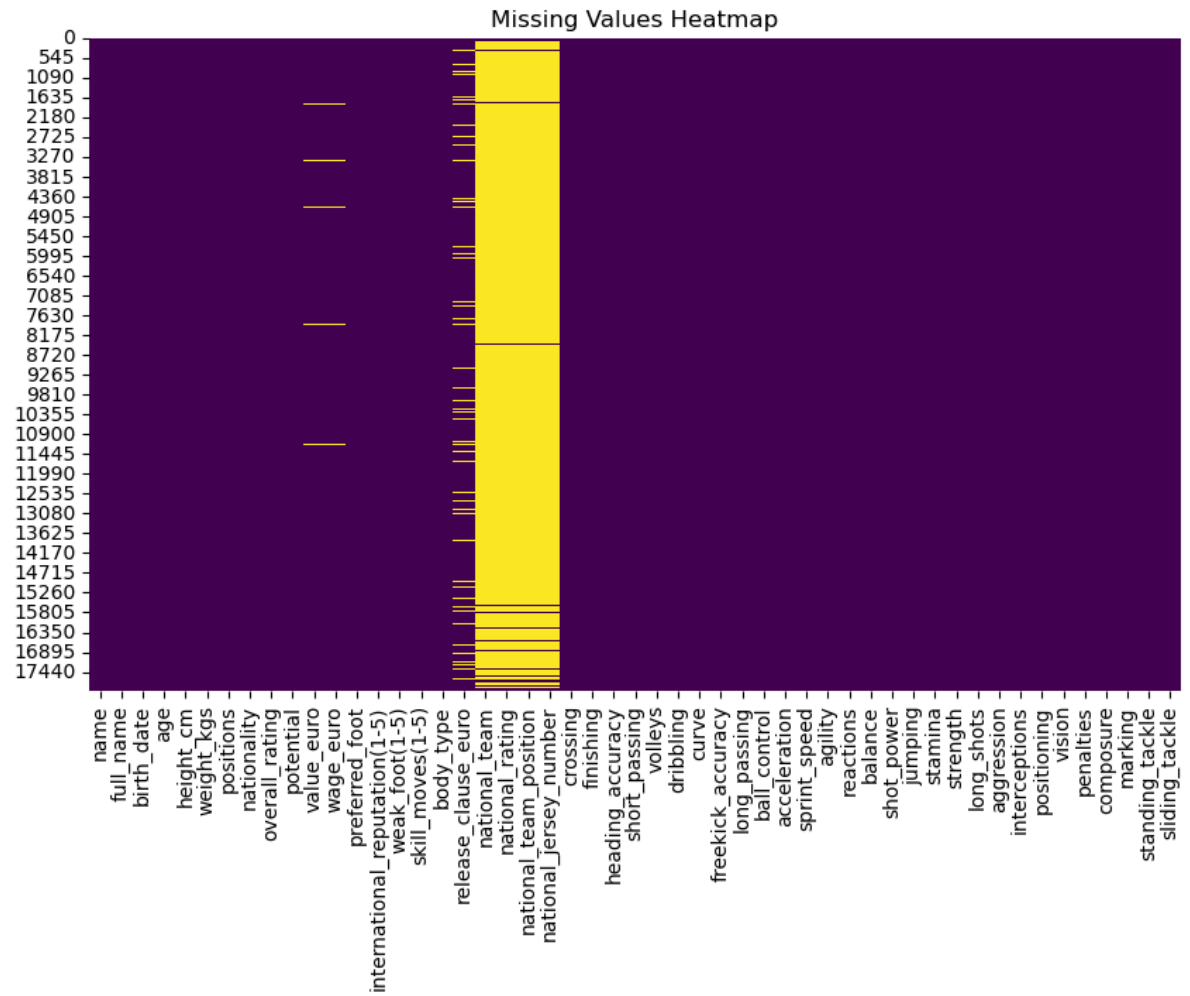
```
In [5]: # Check for missing values
missing_values = fifa_data.isnull().sum().sort_values(ascending=False)
print("\nMissing Values:")
print(missing_values[missing_values > 0])
```

Missing Values:

national_jersey_number	17097
national_team_position	17097
national_rating	17097
national_team	17097
release_clause_euro	1837
value_euro	255
wage_euro	246

dtype: int64

```
In [6]: # Visualize missing data
plt.figure(figsize=(10, 6))
sns.heatmap(fifa_data.isnull(), cbar=False, cmap="viridis")
plt.title("Missing Values Heatmap")
plt.show()
```



```
In [7]: # Handle missing values
cleaned_data = fifa_data.dropna(subset=["overall_rating", "potential", "value_euro"])
cleaned_data['release_clause_euro'].fillna(cleaned_data['release_clause_euro'].median(), inplace=True)
cleaned_data['national_rating'].fillna(0, inplace=True)
cleaned_data['national_jersey_number'].fillna(0, inplace=True)
```

C:\Users\mkhan\AppData\Local\Temp\ipykernel_27028\126606181.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
cleaned_data['release_clause_euro'].fillna(cleaned_data['release_clause_euro'].median(), inplace=True)
```

C:\Users\mkhan\AppData\Local\Temp\ipykernel_27028\126606181.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
cleaned_data['national_rating'].fillna(0, inplace=True)
```

C:\Users\mkhan\AppData\Local\Temp\ipykernel_27028\126606181.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

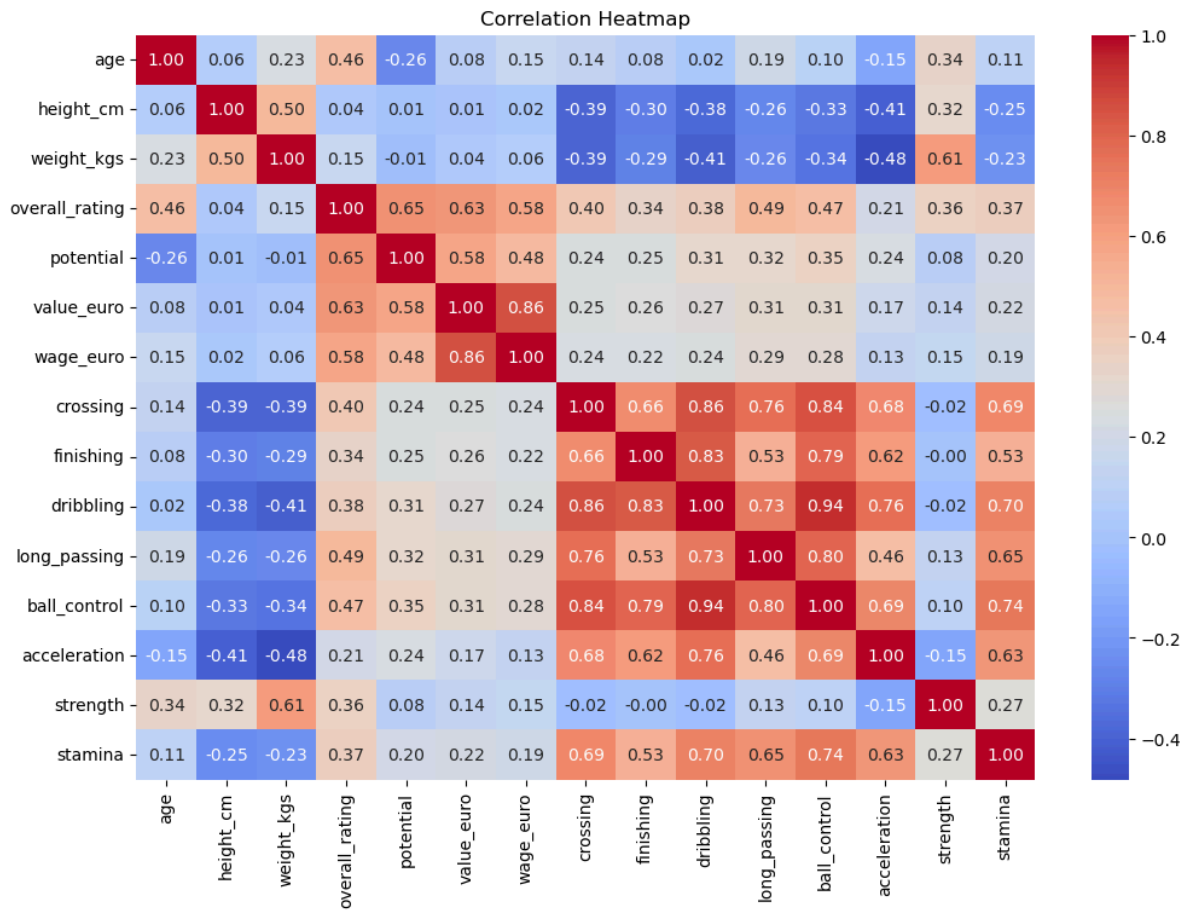
```
cleaned_data['national_jersey_number'].fillna(0, inplace=True)
```

```
In [8]: # Encode categorical variables
encoded_data = pd.get_dummies(cleaned_data, columns=['preferred_foot', 'body_type'])
```

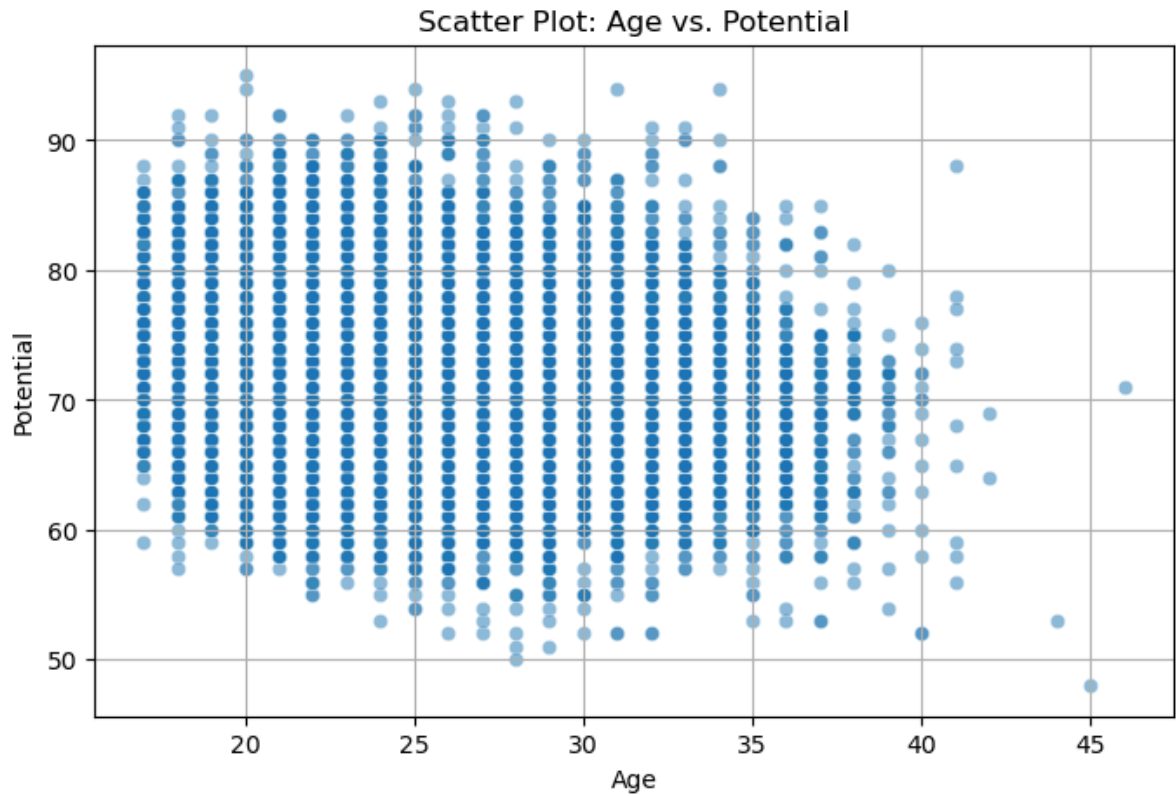
```
In [9]: # Feature correlations
features = [
    "age", "height_cm", "weight_kgs", "overall_rating", "potential", "value_euro",
    "crossing", "finishing", "dribbling", "long_passing", "ball_control", "acceleration",
    "strength", "stamina"
]
correlation_matrix = cleaned_data[features].corr()
```



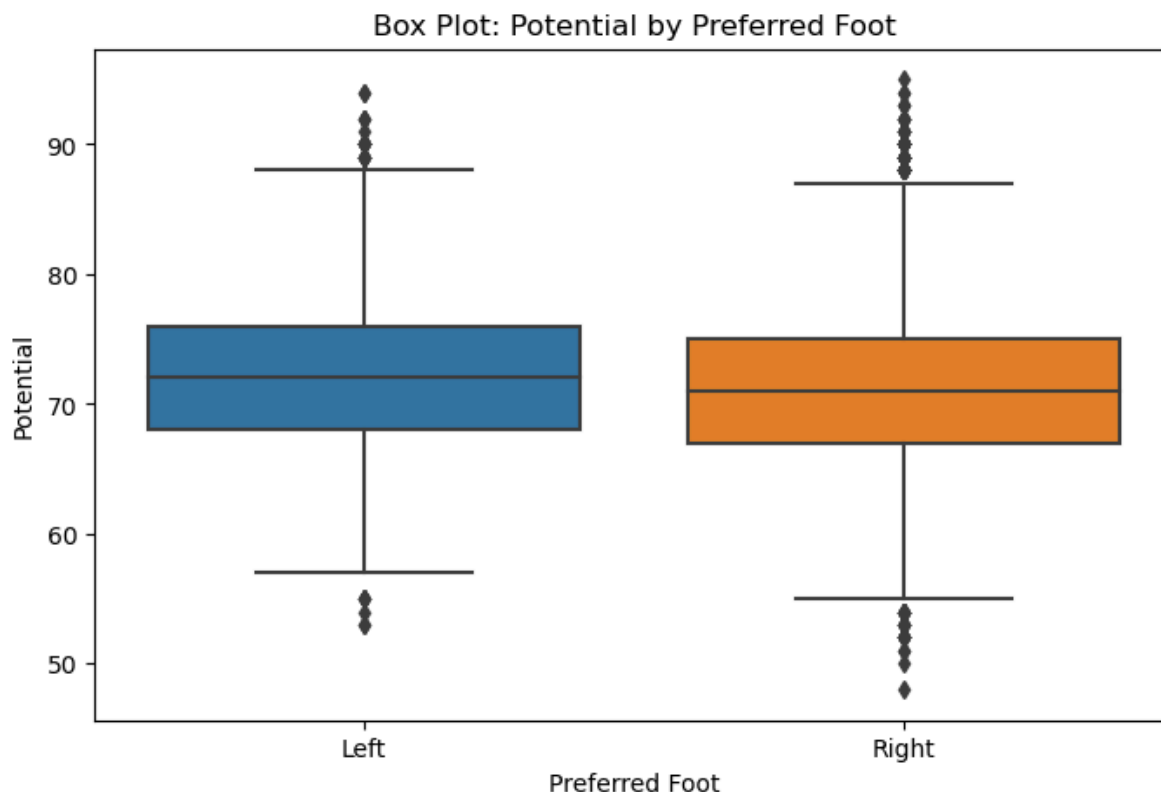
```
In [10]: plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()
```



```
In [28]: # Scatter plot: Age vs. Potential
plt.figure(figsize=(8, 5))
sns.scatterplot(data=fifa_data, x='age', y='potential', alpha=0.5)
plt.title("Scatter Plot: Age vs. Potential")
plt.xlabel("Age")
plt.ylabel("Potential")
plt.grid(True)
plt.show()
```



```
In [29]: # Box plot: Potential by Preferred Foot
plt.figure(figsize=(8, 5))
sns.boxplot(data=fifa_data, x='preferred_foot', y='potential')
plt.title("Box Plot: Potential by Preferred Foot")
plt.xlabel("Preferred Foot")
plt.ylabel("Potential")
plt.show()
```



Machine Learning Section

```
In [11]: # Select features and target
model_features = [
    "age", "height_cm", "weight_kgs", "value_euro", "wage_euro", "crossing",
    "finishing", "dribbling", "long_passing", "ball_control", "acceleration",
    "strength", "stamina", "preferred_foot_Right", "body_type_Lean"
]
target = "potential"
```

```
In [12]: X = encoded_data[model_features]
y = encoded_data[target]
```

```
In [13]: # Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [14]: # Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
```

```
In [15]: # Train Random Forest Regressor
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
```

```
Out[15]:
RandomForestRegressor
RandomForestRegressor(random_state=42)
```

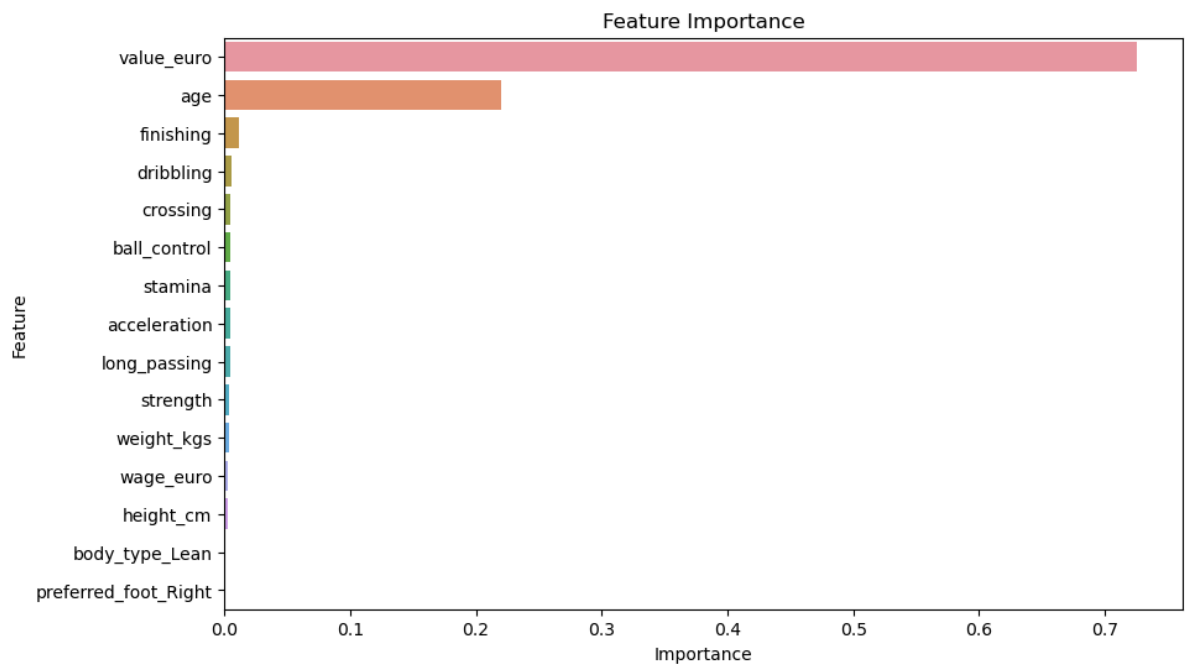
```
In [16]: # Evaluate model
y_pred = rf_model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Random Forest MSE: {mse:.2f}")
print(f"Random Forest R2 Score: {r2:.2f}")
```

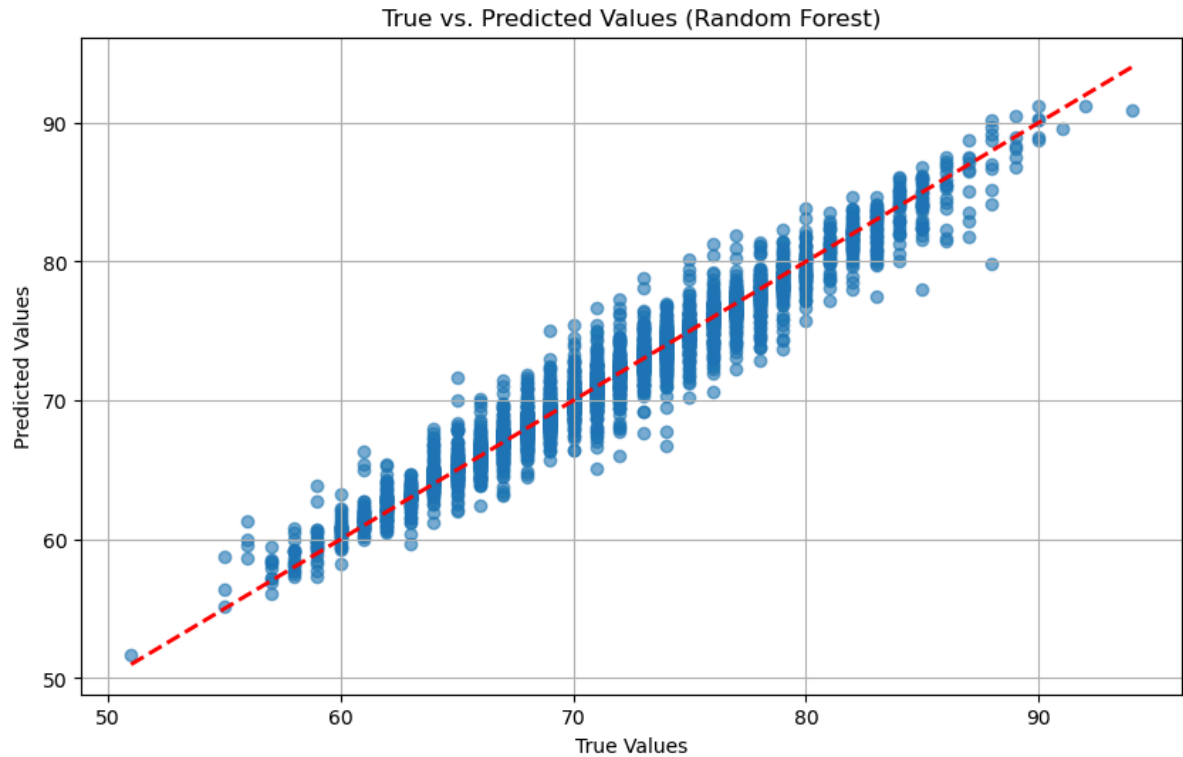
```
Random Forest MSE: 1.81
Random Forest R2 Score: 0.95
```

```
In [17]: # Feature importance
feature_importances = pd.DataFrame({
    'Feature': model_features,
    'Importance': rf_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

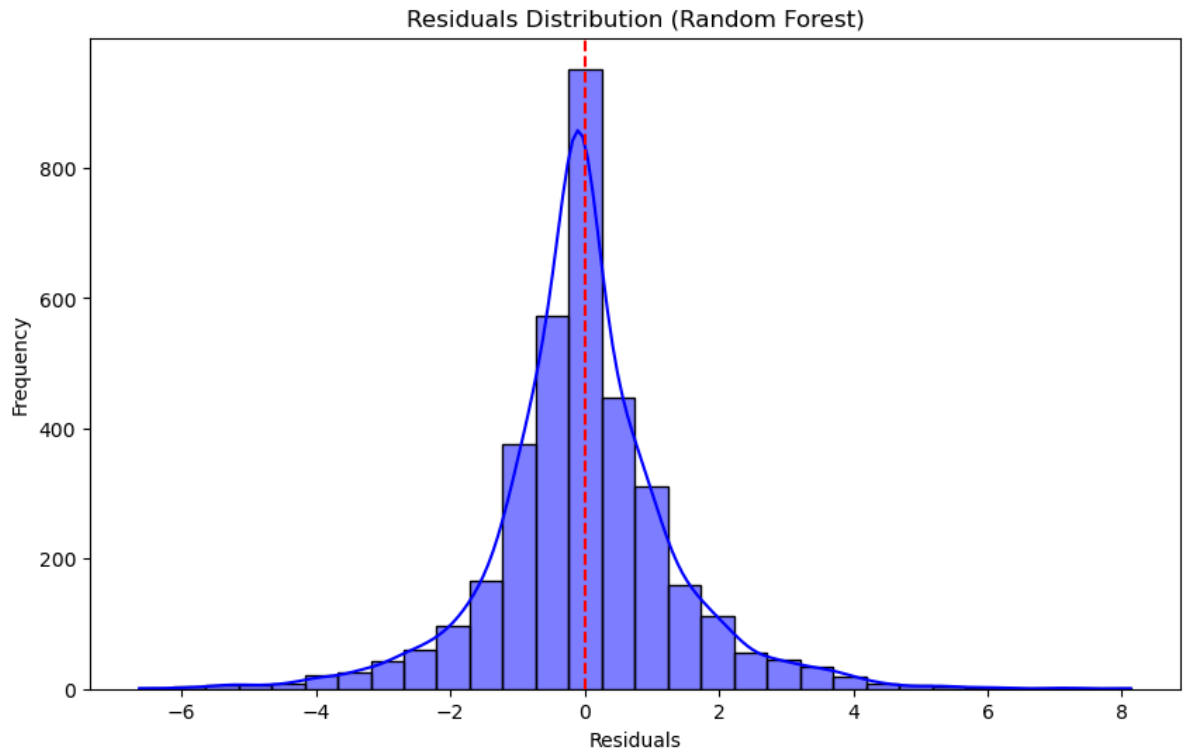
# Visualize feature importance
plt.figure(figsize=(10, 6))
sns.barplot(x=feature_importances['Importance'], y=feature_importances['Feature'])
plt.title("Feature Importance")
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.show()
```



```
In [23]: # True vs. Predicted Values Plot
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red')
plt.title("True vs. Predicted Values (Random Forest)")
plt.xlabel("True Values")
plt.ylabel("Predicted Values")
plt.grid(True)
plt.show()
```



```
In [24]: # Residual Plot
residuals = y_test - y_pred
plt.figure(figsize=(10, 6))
sns.histplot(residuals, bins=30, kde=True, color='blue')
plt.title("Residuals Distribution (Random Forest)")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.axvline(0, color='red', linestyle='--')
plt.show()
```

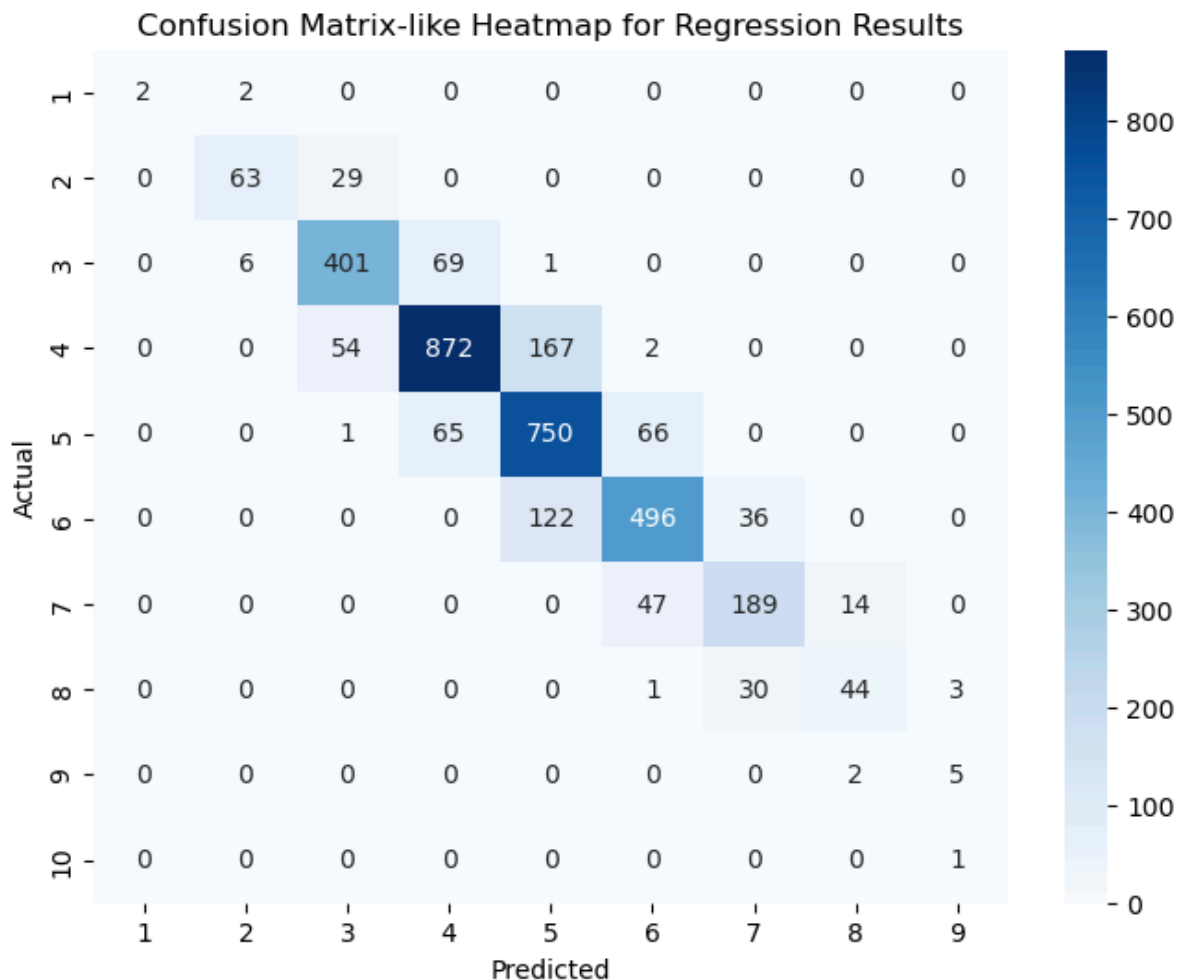


```
In [30]: # Confusion matrix-like visualization for regression results
from sklearn.metrics import mean_squared_error
import numpy as np

# Categorize predictions and true values into ranges
bins = np.linspace(min(y_test), max(y_test), 10)
y_test_bins = np.digitize(y_test, bins)
y_pred_bins = np.digitize(y_pred, bins)

confusion_matrix = pd.crosstab(y_test_bins, y_pred_bins, rownames=['Actual'], columns=['Predicted'], margins=True)

plt.figure(figsize=(8, 6))
sns.heatmap(confusion_matrix, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix-like Heatmap for Regression Results")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



In []:


Deep Learning Section



```
In [18]: # Define the neural network
dl_model = Sequential([
    Dense(128, input_dim=X_train.shape[1], activation='relu'),
    Dropout(0.3),
    Dense(64, activation='relu'),
    Dropout(0.3),
    Dense(32, activation='relu'),
    Dense(1, activation='linear') # Regression output
])
```


C:\Users\mkhan\AppData\Roaming\Python\Python311\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)


```
In [19]: # Compile the model
dl_model.compile(optimizer='adam', loss='mse', metrics=['mae'])
```


```
In [20]: # Train the model
history = dl_model.fit(X_train, y_train, validation_split=0.2, epochs=50, batch
```


Epoch 1/50
354/354  2s 2ms/step - loss: 2360.8167 - mae: 39.9123 - val_loss: 61.2291 - val_mae: 5.7972


Epoch 2/50
354/354  0s 1ms/step - loss: 133.2749 - mae: 9.1570 - val_loss: 29.1820 - val_mae: 4.0614


Epoch 3/50
354/354  1s 1ms/step - loss: 109.2854 - mae: 8.3019 - val_loss: 23.2581 - val_mae: 3.6729


Epoch 4/50
354/354  0s 1ms/step - loss: 89.4665 - mae: 7.5339 - val_loss: 19.0750 - val_mae: 3.2351


Epoch 5/50
354/354  0s 1ms/step - loss: 79.6058 - mae: 7.0991 - val_loss: 24.4012 - val_mae: 4.0042


Epoch 6/50
354/354  1s 1ms/step - loss: 79.4471 - mae: 7.0914 - val_loss: 21.7368 - val_mae: 3.5744


Epoch 7/50
354/354  0s 1ms/step - loss: 74.7454 - mae: 6.9061 - val_loss: 16.2896 - val_mae: 3.1400


Epoch 8/50
354/354  0s 1ms/step - loss: 73.5921 - mae: 6.7783 - val_loss: 14.9658 - val_mae: 2.8747


Epoch 9/50
354/354  0s 1ms/step - loss: 68.1473 - mae: 6.5480 - val_loss: 13.2049 - val_mae: 2.7067


Epoch 10/50
354/354  1s 1ms/step - loss: 63.1669 - mae: 6.2904 - val_loss: 14.5305 - val_mae: 2.9179


Epoch 11/50
354/354  0s 1ms/step - loss: 60.4002 - mae: 6.1915 - val_loss: 18.8306 - val_mae: 3.4872


Epoch 12/50
354/354  1s 1ms/step - loss: 58.5381 - mae: 6.0857 - val_loss: 13.0532 - val_mae: 2.7736


Epoch 13/50
354/354  1s 1ms/step - loss: 57.8573 - mae: 6.0327 - val_loss: 13.3369 - val_mae: 2.8023


Epoch 14/50
354/354  1s 1ms/step - loss: 55.0884 - mae: 5.8994 - val_loss: 13.3007 - val_mae: 2.7793


Epoch 15/50
354/354  1s 2ms/step - loss: 49.4275 - mae: 5.5601 - val_loss: 16.2631 - val_mae: 3.2908


Epoch 16/50
354/354  1s 1ms/step - loss: 48.2945 - mae: 5.5408 - val_loss: 13.2673 - val_mae: 2.8274


Epoch 17/50
354/354  1s 1ms/step - loss: 44.4011 - mae: 5.3063 - val_loss: 13.5689 - val_mae: 2.8373


Epoch 18/50
354/354  1s 2ms/step - loss: 43.4567 - mae: 5.1965 - val_loss: 12.9180 - val_mae: 2.8198


Epoch 19/50
354/354  2s 4ms/step - loss: 41.6199 - mae: 5.0866 - val_loss: 11.4982 - val_mae: 2.6297


Epoch 20/50
354/354  1s 2ms/step - loss: 38.5464 - mae: 4.9256 - val_loss: 12.7631 - val_mae: 2.7909


Epoch 21/50
354/354  1s 2ms/step - loss: 36.6037 - mae: 4.7633 - val_loss: 14.0959 - val_mae: 2.9133


Epoch 22/50
354/354  1s 1ms/step - loss: 33.6952 - mae: 4.5895 - val_loss: 25.4520 - val_mae: 4.1988


Epoch 23/50
354/354  1s 2ms/step - loss: 31.8507 - mae: 4.4386 - val_loss: 43.4915 - val_mae: 5.7719


Epoch 24/50
354/354  1s 2ms/step - loss: 29.1082 - mae: 4.2363 - val_loss: 32.0432 - val_mae: 4.8280


Epoch 25/50
354/354  1s 2ms/step - loss: 26.9354 - mae: 4.0702 - val_loss: 32.8055 - val_mae: 4.9889


Epoch 26/50
354/354  1s 2ms/step - loss: 24.0457 - mae: 3.8290 - val_loss: 39.8793 - val_mae: 5.6049


Epoch 27/50
354/354  1s 2ms/step - loss: 22.8041 - mae: 3.6860 - val_loss: 41.4118 - val_mae: 5.7254


Epoch 28/50
354/354  1s 2ms/step - loss: 21.7078 - mae: 3.6059 - val_loss: 52.2183 - val_mae: 6.6180


Epoch 29/50
354/354  1s 2ms/step - loss: 21.5858 - mae: 3.5603 - val_loss: 40.3296 - val_mae: 5.7671


Epoch 30/50
354/354  0s 1ms/step - loss: 20.0841 - mae: 3.4516 - val_loss: 56.0655 - val_mae: 6.9013


Epoch 31/50
354/354  0s 1ms/step - loss: 17.9723 - mae: 3.2986 - val_loss: 44.1900 - val_mae: 6.0435

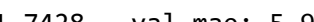
Epoch 32/50
354/354  1s 2ms/step - loss: 17.1947 - mae: 3.1773 - val_loss: 60.9792 - val_mae: 7.2623

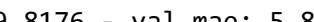
Epoch 33/50
354/354  1s 3ms/step - loss: 16.1487 - mae: 3.0562 - val_loss: 58.4195 - val_mae: 7.1509













Epoch 34/50
354/354  1s 2ms/step - loss: 15.3841 - mae: 3.0173 - val_loss: 50.0336 - val_mae: 6.5905

Epoch 35/50
354/354  1s 2ms/step - loss: 13.7143 - mae: 2.8471 - val_loss: 40.9052 - val_mae: 5.8958

Epoch 36/50
354/354  1s 2ms/step - loss: 13.5447 - mae: 2.8277 - val_loss: 42.5010 - val_mae: 6.0278

Epoch 37/50
354/354  1s 2ms/step - loss: 13.4804 - mae: 2.8099 - val_loss: 41.7428 - val_mae: 5.9879

Epoch 38/50
354/354  1s 1ms/step - loss: 12.5287 - mae: 2.7023 - val_loss: 39.8176 - val_mae: 5.8396

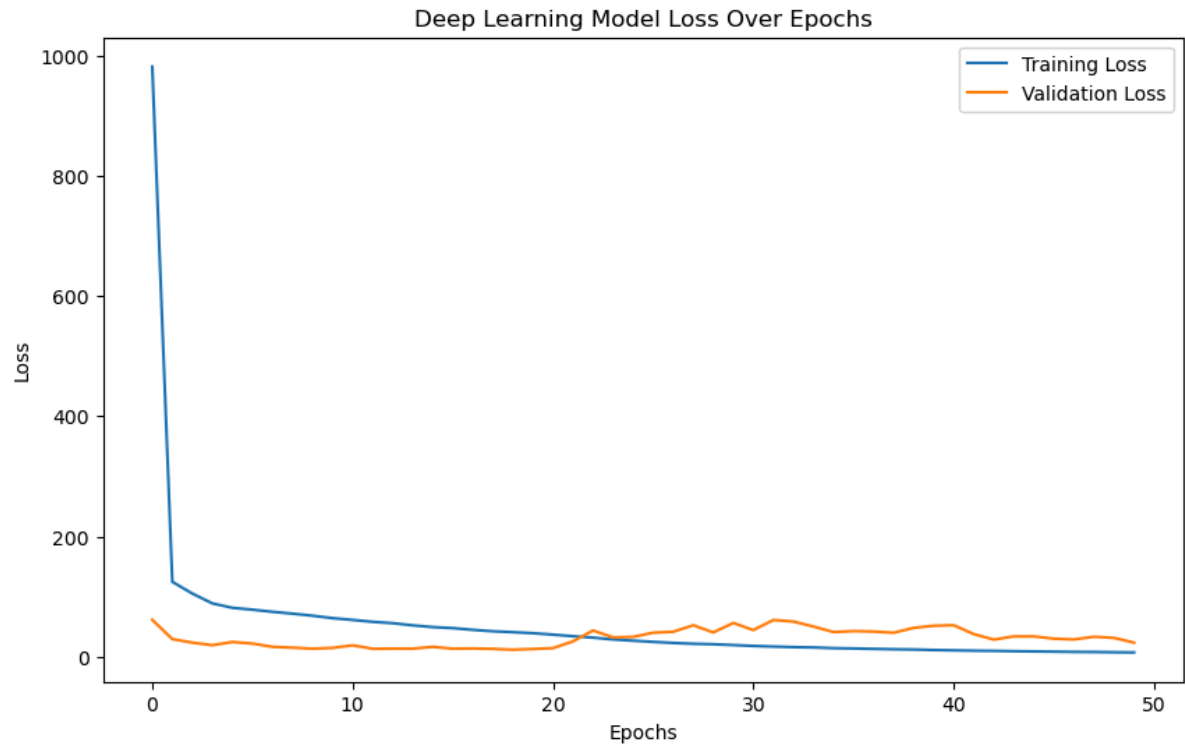
Epoch 39/50
 354/354  0s 1ms/step - loss: 12.1136 - mae: 2.6617 - val_loss: 47.8619 - val_mae: 6.4807
 Epoch 40/50
 354/354  1s 1ms/step - loss: 11.0504 - mae: 2.5363 - val_loss: 51.3326 - val_mae: 6.7263
 Epoch 41/50
 354/354  0s 1ms/step - loss: 10.9035 - mae: 2.5096 - val_loss: 52.3017 - val_mae: 6.7635
 Epoch 42/50
 354/354  1s 1ms/step - loss: 10.3730 - mae: 2.4532 - val_loss: 37.5171 - val_mae: 5.6148
 Epoch 43/50
 354/354  0s 1ms/step - loss: 9.9946 - mae: 2.3858 - val_loss: 28.3320 - val_mae: 4.8971
 Epoch 44/50
 354/354  0s 1ms/step - loss: 8.9855 - mae: 2.2788 - val_loss: 33.6376 - val_mae: 5.4146
 Epoch 45/50
 354/354  0s 1ms/step - loss: 8.7830 - mae: 2.2327 - val_loss: 33.6458 - val_mae: 5.3894
 Epoch 46/50
 354/354  0s 1ms/step - loss: 8.3171 - mae: 2.1847 - val_loss: 29.8541 - val_mae: 5.0626
 Epoch 47/50
 354/354  0s 1ms/step - loss: 7.9018 - mae: 2.1286 - val_loss: 28.5985 - val_mae: 4.8901
 Epoch 48/50
 354/354  1s 1ms/step - loss: 7.4501 - mae: 2.0468 - val_loss: 33.0337 - val_mae: 5.3221
 Epoch 49/50
 354/354  0s 1ms/step - loss: 7.3509 - mae: 2.0319 - val_loss: 31.2201 - val_mae: 5.1457
 Epoch 50/50
 354/354  0s 1ms/step - loss: 7.1031 - mae: 2.0178 - val_loss: 23.2142 - val_mae: 4.3503

```
In [21]: # Evaluate on test set
dl_results = dl_model.evaluate(X_test, y_test, verbose=0)

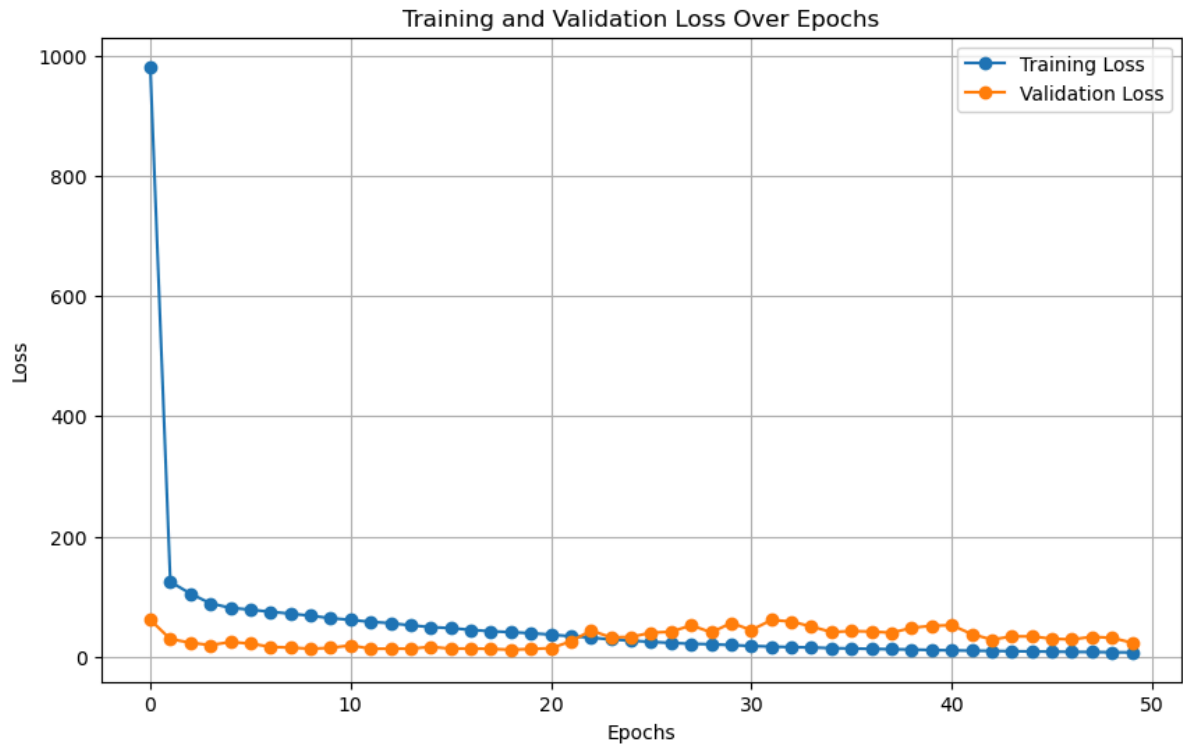
print(f"Deep Learning Test Loss: {dl_results[0]:.2f}")
print(f"Deep Learning Test MAE: {dl_results[1]:.2f}")
```

Deep Learning Test Loss: 22.63
 Deep Learning Test MAE: 4.27

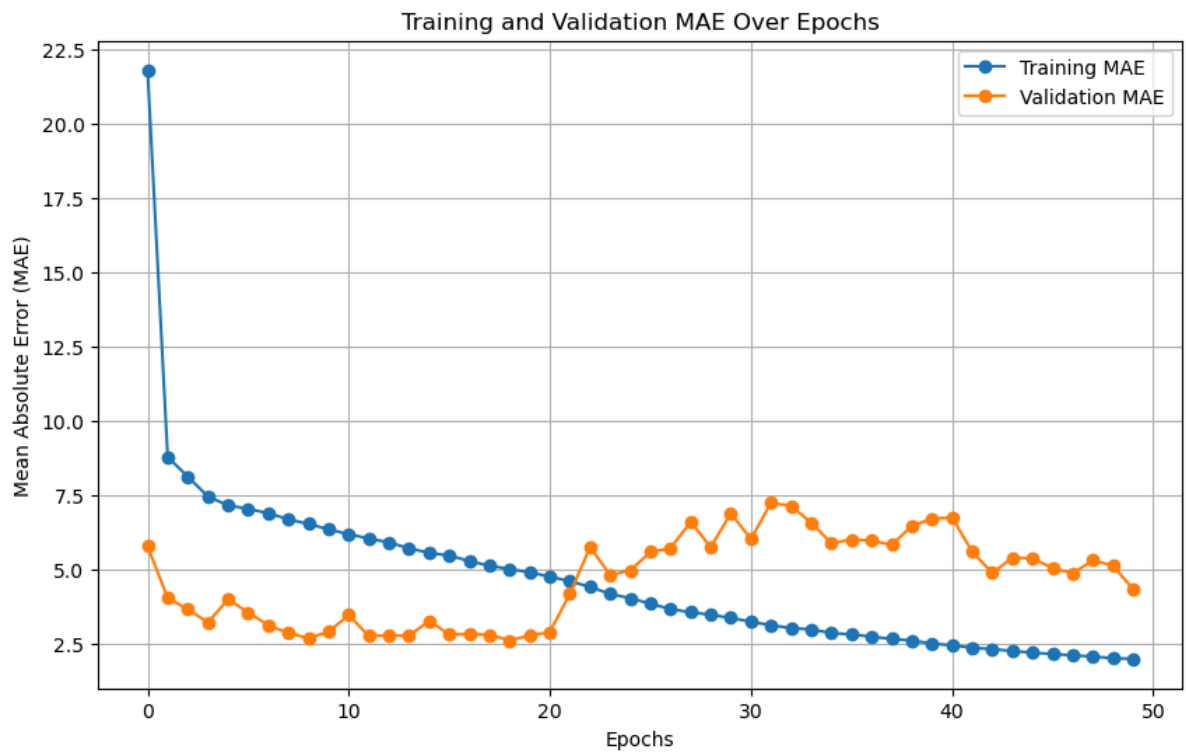
```
In [22]: # Visualize training history
plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title("Deep Learning Model Loss Over Epochs")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



```
In [25]: # Plot Training and Validation Loss
plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'], label='Training Loss', marker='o')
plt.plot(history.history['val_loss'], label='Validation Loss', marker='o')
plt.title("Training and Validation Loss Over Epochs")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()
```



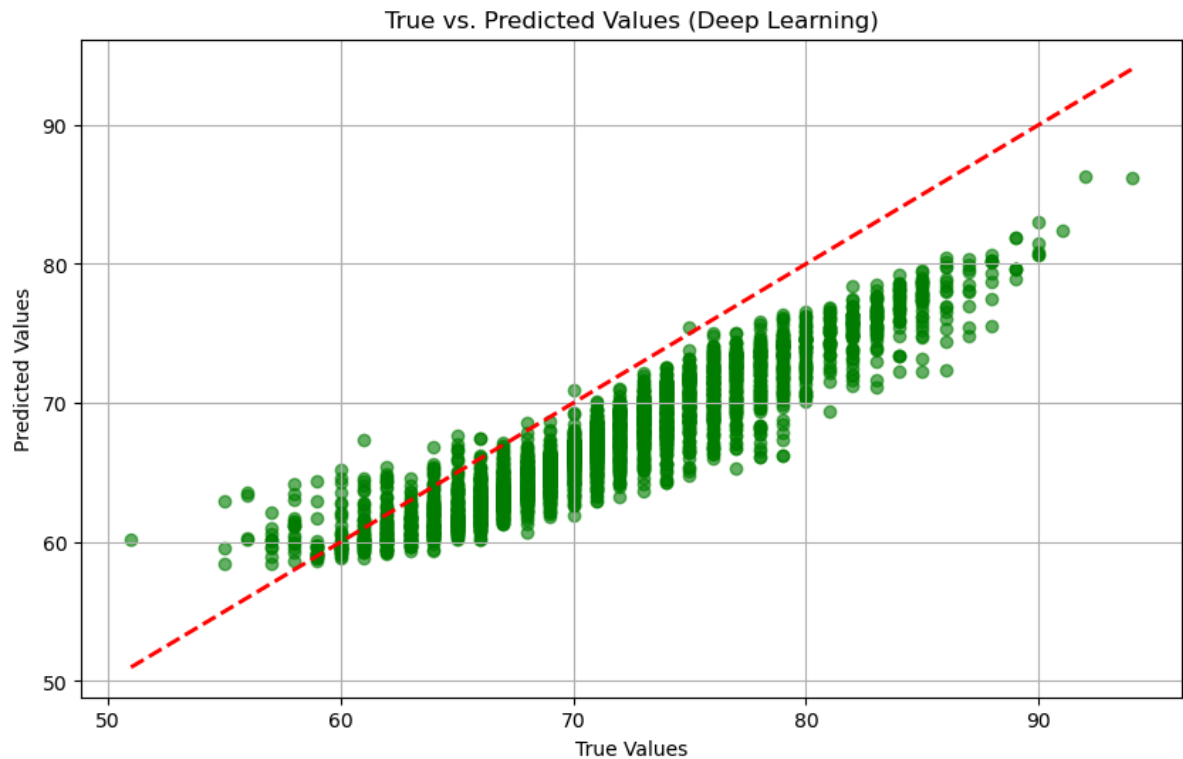
```
In [26]: # Plot Training and Validation MAE
plt.figure(figsize=(10, 6))
plt.plot(history.history['mae'], label='Training MAE', marker='o')
plt.plot(history.history['val_mae'], label='Validation MAE', marker='o')
plt.title("Training and Validation MAE Over Epochs")
plt.xlabel("Epochs")
plt.ylabel("Mean Absolute Error (MAE)")
plt.legend()
plt.grid(True)
plt.show()
```




```
In [27]: # True vs. Predicted Values Scatter Plot
y_dl_pred = dl_model.predict(X_test).flatten()

plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_dl_pred, alpha=0.6, color='green')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red')
plt.title("True vs. Predicted Values (Deep Learning)")
plt.xlabel("True Values")
plt.ylabel("Predicted Values")
plt.grid(True)
plt.show()
```


111/111 ————— 0s 1ms/step





```
In [31]: from tensorflow.keras.callbacks import TensorBoard
import datetime


# TensorBoard setup
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = TensorBoard(log_dir=log_dir, histogram_freq=1)


# Model training with TensorBoard callback
history = dl_model.fit(
    X_train, y_train,
    validation_split=0.2,
    epochs=50,
    batch_size=32,
    verbose=1,
    callbacks=[tensorboard_callback]
)
```


Epoch 1/50
354/354  1s 2ms/step - loss: 7.0558 - mae: 1.9912 - val_loss: 21.8121 - val_mae: 4.2511


Epoch 2/50
354/354  1s 2ms/step - loss: 6.8658 - mae: 1.9418 - val_loss: 19.9683 - val_mae: 4.0161


Epoch 3/50
354/354  1s 2ms/step - loss: 6.4287 - mae: 1.8858 - val_loss: 24.0675 - val_mae: 4.4614


Epoch 4/50
354/354  1s 2ms/step - loss: 6.2698 - mae: 1.8626 - val_loss: 18.1380 - val_mae: 3.8127


Epoch 5/50
354/354  1s 2ms/step - loss: 5.9702 - mae: 1.8067 - val_loss: 16.5177 - val_mae: 3.6317


Epoch 6/50
354/354  1s 2ms/step - loss: 5.7988 - mae: 1.7949 - val_loss: 17.0902 - val_mae: 3.7087


Epoch 7/50
354/354  1s 2ms/step - loss: 5.9902 - mae: 1.8382 - val_loss: 17.6852 - val_mae: 3.8046


Epoch 8/50
354/354  1s 2ms/step - loss: 5.6135 - mae: 1.7671 - val_loss: 17.5584 - val_mae: 3.7491


Epoch 9/50
354/354  1s 2ms/step - loss: 5.5161 - mae: 1.7332 - val_loss: 17.3618 - val_mae: 3.7516


Epoch 10/50
354/354  1s 2ms/step - loss: 5.5013 - mae: 1.7441 - val_loss: 15.7107 - val_mae: 3.4837


Epoch 11/50
354/354  1s 2ms/step - loss: 5.1521 - mae: 1.6864 - val_loss: 17.4275 - val_mae: 3.7205


Epoch 12/50
354/354  1s 2ms/step - loss: 5.2380 - mae: 1.6906 - val_loss: 14.4171 - val_mae: 3.3757


Epoch 13/50
354/354  1s 2ms/step - loss: 5.0566 - mae: 1.6641 - val_loss: 14.5546 - val_mae: 3.4122


Epoch 14/50
354/354  1s 2ms/step - loss: 5.0519 - mae: 1.6630 - val_loss: 17.3743 - val_mae: 3.7308




















Epoch 15/50
354/354  1s 2ms/step - loss: 4.8877 - mae: 1.6400 - val_loss: 15.8087 - val_mae: 3.5656

Epoch 16/50
354/354  1s 2ms/step - loss: 4.8021 - mae: 1.6180 - val_loss: 19.2609 - val_mae: 3.9886


Epoch 17/50
354/354  1s 2ms/step - loss: 4.7583 - mae: 1.6109 - val_loss: 17.3000 - val_mae: 3.7671

Epoch 18/50
354/354  1s 2ms/step - loss: 4.7592 - mae: 1.6062 - val_loss: 13.7713 - val_mae: 3.2916


Epoch 19/50
354/354  1s 2ms/step - loss: 4.6130 - mae: 1.5963 - val_loss: 14.1414 - val_mae: 3.3210

Epoch 20/50
354/354  1s 2ms/step - loss: 4.5580 - mae: 1.5847 - val_loss: 15.9713 - val_mae: 3.5325
Epoch 21/50
354/354  1s 2ms/step - loss: 4.4756 - mae: 1.5552 - val_loss: 15.5634 - val_mae: 3.4834
Epoch 22/50
354/354  1s 2ms/step - loss: 4.4192 - mae: 1.5609 - val_loss: 13.6270 - val_mae: 3.2891
Epoch 23/50
354/354  1s 2ms/step - loss: 4.3514 - mae: 1.5314 - val_loss: 16.1044 - val_mae: 3.6045
Epoch 24/50
354/354  1s 2ms/step - loss: 4.6010 - mae: 1.5634 - val_loss: 14.0938 - val_mae: 3.3067
Epoch 25/50
354/354  1s 2ms/step - loss: 4.3295 - mae: 1.5378 - val_loss: 11.5513 - val_mae: 2.9376
Epoch 26/50
354/354  1s 2ms/step - loss: 4.2732 - mae: 1.5205 - val_loss: 11.7098 - val_mae: 3.0031
Epoch 27/50
354/354  1s 2ms/step - loss: 4.1749 - mae: 1.4920 - val_loss: 13.4141 - val_mae: 3.2194
Epoch 28/50
354/354  1s 2ms/step - loss: 4.2798 - mae: 1.5127 - val_loss: 13.0060 - val_mae: 3.2156
Epoch 29/50
354/354  1s 2ms/step - loss: 4.2907 - mae: 1.5295 - val_loss: 13.7055 - val_mae: 3.2363
Epoch 30/50
354/354  1s 2ms/step - loss: 4.1720 - mae: 1.4965 - val_loss: 12.5416 - val_mae: 3.1174
Epoch 31/50
354/354  1s 2ms/step - loss: 4.0325 - mae: 1.4821 - val_loss: 10.5649 - val_mae: 2.9054
Epoch 32/50
354/354  1s 2ms/step - loss: 3.9882 - mae: 1.4671 - val_loss: 11.7065 - val_mae: 2.9864
Epoch 33/50
354/354  1s 2ms/step - loss: 4.0279 - mae: 1.4866 - val_loss: 12.2868 - val_mae: 3.1479
Epoch 34/50
354/354  1s 2ms/step - loss: 4.0007 - mae: 1.4737 - val_loss: 12.2819 - val_mae: 3.1129
Epoch 35/50
354/354  1s 2ms/step - loss: 4.1805 - mae: 1.5045 - val_loss: 10.3223 - val_mae: 2.8056
Epoch 36/50
354/354  1s 2ms/step - loss: 3.9037 - mae: 1.4542 - val_loss: 13.2681 - val_mae: 3.1773
Epoch 37/50
354/354  1s 2ms/step - loss: 3.9669 - mae: 1.4705 - val_loss: 14.6843 - val_mae: 3.4455
Epoch 38/50
354/354  1s 2ms/step - loss: 4.0510 - mae: 1.4617 - val_loss: 10.0642 - val_mae: 2.7988


Epoch 39/50

354/354  1s 2ms/step - loss: 3.7931 - mae: 1.4254 - val_loss: 11.2842 - val_mae: 2.9759


Epoch 40/50

354/354  1s 2ms/step - loss: 3.7446 - mae: 1.4186 - val_loss: 14.4052 - val_mae: 3.3422


Epoch 41/50

354/354  1s 2ms/step - loss: 3.9305 - mae: 1.4456 - val_loss: 12.7288 - val_mae: 3.1599


Epoch 42/50

354/354  1s 2ms/step - loss: 3.5872 - mae: 1.3996 - val_loss: 12.1998 - val_mae: 3.0624


Epoch 43/50

354/354  1s 2ms/step - loss: 4.1423 - mae: 1.4687 - val_loss: 13.0107 - val_mae: 3.1933


Epoch 44/50

354/354  1s 2ms/step - loss: 3.9138 - mae: 1.4234 - val_loss: 10.9625 - val_mae: 2.8861


Epoch 45/50

354/354  1s 2ms/step - loss: 3.7473 - mae: 1.4266 - val_loss: 13.0028 - val_mae: 3.1772


Epoch 46/50

354/354  1s 2ms/step - loss: 3.9098 - mae: 1.4357 - val_loss: 10.2345 - val_mae: 2.7965


Epoch 47/50

354/354  1s 2ms/step - loss: 3.6212 - mae: 1.3955 - val_loss: 10.8252 - val_mae: 2.8630


Epoch 48/50

354/354  1s 2ms/step - loss: 3.8664 - mae: 1.4526 - val_loss: 10.8232 - val_mae: 2.8617

Epoch 49/50

354/354  1s 2ms/step - loss: 3.8166 - mae: 1.4182 - val_loss: 12.5546 - val_mae: 3.1572

Epoch 50/50

354/354  1s 2ms/step - loss: 3.7474 - mae: 1.4122 - val_loss: 12.2438 - val_mae: 3.0675

```
In [ ]: from sklearn.model_selection import GridSearchCV

# Define parameter grid
param_grid = {
    'n_estimators': [50, 100, 150],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5, 10]
}

# Grid search
grid_search = GridSearchCV(estimator=RandomForestRegressor(random_state=42), param_grid=param_grid)
grid_search.fit(X_train, y_train)

# Best parameters
print("Best Parameters:", grid_search.best_params_)
print("Best R2 Score:", grid_search.best_score_)
```

Fitting 3 folds for each of 27 candidates, totalling 81 fits

```
[CV] END .max_depth=10, min_samples_split=2, n_estimators=50; total time=
3.2s
[CV] END .max_depth=10, min_samples_split=2, n_estimators=50; total time=
3.2s
[CV] END .max_depth=10, min_samples_split=2, n_estimators=50; total time=
3.2s
[CV] END max_depth=10, min_samples_split=2, n_estimators=100; total time=
6.9s
[CV] END max_depth=10, min_samples_split=2, n_estimators=100; total time=
7.3s
[CV] END max_depth=10, min_samples_split=2, n_estimators=100; total time=
7.6s
[CV] END max_depth=10, min_samples_split=2, n_estimators=150; total time= 1
1.1s
[CV] END max_depth=10, min_samples_split=2, n_estimators=150; total time= 1
1.0s
[CV] END max_depth=10, min_samples_split=2, n_estimators=150; total time= 1
1.6s
[CV] END .max_depth=10, min_samples_split=5, n_estimators=50; total time=
3.7s
[CV] END .max_depth=10, min_samples_split=5, n_estimators=50; total time=
3.7s
[CV] END .max_depth=10, min_samples_split=5, n_estimators=50; total time=
3.7s
[CV] END max_depth=10, min_samples_split=5, n_estimators=100; total time=
7.8s
[CV] END max_depth=10, min_samples_split=5, n_estimators=100; total time=
7.8s
[CV] END max_depth=10, min_samples_split=5, n_estimators=100; total time=
6.6s
[CV] END max_depth=10, min_samples_split=5, n_estimators=150; total time= 1
1.3s
[CV] END max_depth=10, min_samples_split=5, n_estimators=150; total time= 1
0.4s
[CV] END max_depth=10, min_samples_split=5, n_estimators=150; total time= 1
0.0s
[CV] END max_depth=10, min_samples_split=10, n_estimators=50; total time=
3.1s
[CV] END max_depth=10, min_samples_split=10, n_estimators=50; total time=
3.4s
[CV] END max_depth=10, min_samples_split=10, n_estimators=50; total time=
3.0s
[CV] END max_depth=10, min_samples_split=10, n_estimators=100; total time=
6.2s
[CV] END max_depth=10, min_samples_split=10, n_estimators=100; total time=
7.5s
[CV] END max_depth=10, min_samples_split=10, n_estimators=100; total time=
6.7s
[CV] END max_depth=10, min_samples_split=10, n_estimators=150; total time= 1
0.2s
[CV] END max_depth=10, min_samples_split=10, n_estimators=150; total time= 1
0.3s
[CV] END max_depth=10, min_samples_split=10, n_estimators=150; total time=
9.7s
[CV] END .max_depth=20, min_samples_split=2, n_estimators=50; total time=
6.2s
```

```
[CV] END .max_depth=20, min_samples_split=2, n_estimators=50; total time= 6.1s
[CV] END .max_depth=20, min_samples_split=2, n_estimators=50; total time= 6.0s
[CV] END max_depth=20, min_samples_split=2, n_estimators=100; total time= 12.5s
[CV] END max_depth=20, min_samples_split=2, n_estimators=100; total time= 12.8s
[CV] END max_depth=20, min_samples_split=2, n_estimators=100; total time= 12.5s
[CV] END max_depth=20, min_samples_split=2, n_estimators=150; total time= 18.6s
[CV] END max_depth=20, min_samples_split=2, n_estimators=150; total time= 20.9s
```

Conclusions

In []: EDA revealed strong correlations between potential **and** attributes like value_eu
Machine Learning: Random Forest achieved score of **~0.85**, indicating good predic
Deep Learning: The neural network effectively modeled player potential **with** a n

References

- Academic (if any)
- Online (if any)

In []:

Credits

- If you use and/or adapt your code from existing projects, you must provide links and acknowledge the authors.

This code is based on (if any)

In []:

In []: **# End of Project**